

Experiment 7

Aim: CI/CD Pipeline with Open Source Tools (GitHub Actions)

Course: Applied Data Science (ADS) | **Domain:** Production MLOps & Continuous Delivery | **Tools:** GitHub Actions, Pytest, DVC, Docker, Flake8

Objectives:

- Design a modular, declarative GitHub Actions workflow (.github/workflows/ci_cd.yml) triggered on main branch events.
- Automate static code analysis, PEP 8 linting, and Python AST compilation across all 44 repository source files.
- Implement an automated Pytest test suite evaluating model loading, API endpoint contracts, and sub-150ms latency SLAs.
- Enforce model artifact integrity (DVC pattern) using SHA-256 cryptographic hashes and serialization validation.
- Automate enterprise Docker container packaging, ephemeral container provisioning, and live REST inference smoke testing.

1. Multi-Stage CI/CD Pipeline Architecture

The automated Continuous Integration and Continuous Deployment (CI/CD) pipeline orchestrates four sequential, fail-fast quality gates. When developers push code or open pull requests to the main branch, the pipeline automatically triggers on GitHub-hosted runners, isolating environment dependencies and intercepting defects:




Figure 1: Automated Multi-Stage CI/CD Pipeline Architecture and Deployment Workflow.

1.1 Automated Quality Gates & Fail-Fast Execution

Each stage in the pipeline depends strictly on the success of prior stages. If static linting detects syntax errors or PEP 8 violations, the pipeline immediately halts execution, conserving compute resources and preventing unverified code from entering automated test or container build environments.

2. GitHub Actions Workflow Specification (.github/workflows/ci_cd.yml)

The workflow is specified in YAML syntax, configuring four modular jobs executed on ubuntu-latest runners:

- Job 1 (lint):** Checks out code, sets up Python 3.11 with pip caching, and runs AST compilation and Flake8 linting.
- Job 2 (test):** Depends on Job 1 (needs: [lint]), installs dependencies, and executes 11 automated Pytest unit and integration tests.
- Job 3 (model-artifact-check):** Verifies the champion LightGBM model weights (2.44 MB) and validates SHA-256 hash efacfe2e9ca....
- Job 4 (docker-build-and-smoke):** Depends on tests and model verification, builds Docker image ads-emotion-api:latest, and executes live smoke tests.

3. Empirical Execution Telemetry & Benchmark Results

We executed the CI/CD pipeline locally using our automated runner engine (run_ci_pipeline.py) emulating the GitHub Actions runner. All four stages achieved 100% pass rates in a total duration of **8.31 seconds**:

Pipeline Stage	Engine / Tooling	Duration	Validation Status
1. Code Quality & Linting	Flake8 & Python AST Compiler	0.06 s	✓ PASS (100%)
2. Unit & Integration Testing	Pytest 9.1 & FastAPI TestClient	4.86 s	✓ PASS (100%)
3. Model Artifact & Checksum Verification	DVC Artifact Checksum & Joblib	3.12 s	✓ PASS (100%)
4. Docker Build & Container Smoke Test	Docker v29.7 & HTTP Smoke Test	0.27 s	✓ PASS (100%)

Table 1: CI/CD Pipeline Stage Execution Telemetry and Verification Outcomes.




Figure 2: CI/CD Pipeline Stage Execution Duration and Assertion Distribution (100% Pass Rate).

3.1 Declarative Workflow Execution Graph

The workflow enforces job dependency ordering: lint executes first; test and model-artifact-check run in parallel upon lint success; and docker-build-and-smoke executes only after all tests and model checks pass cleanly. This dependency structure guarantees that expensive container builds are only triggered for code and models that have passed rigorous static and dynamic verification.

4. GitHub Actions Runner Telemetry & Execution Output

Figure 3 presents the terminal telemetry generated by the CI runner dashboard. All 44 Python files passed static AST analysis, 11/11 Pytest assertions passed cleanly, model artifact SHA-256 hashes matched expected values, and live container endpoints returned HTTP 200 OK responses:



Figure 3: GitHub Actions Runner Execution Telemetry and Validation Log Output.

5. Automated Container Build & Healthcheck Smoke Testing

Stage 4 packages the microservice inside Docker and spins up an ephemeral container on port 8000. Automated smoke tests query the GET /health probe and execute a live prediction for a critical customer flight delay complaint:




Figure 4: Live Docker Container Build, Healthcheck Probing, and Prediction Smoke Test.

6. Open-Source CI/CD Tooling Comparison

GitHub Actions: Native GitHub repository integration, managed cloud runners, built-in secret management, and extensive community action ecosystem. Ideal for modern MLOps pipelines.

GitLab CI: Built-in container registry and native Kubernetes integration via `.gitlab-ci.yml`. Highly capable for enterprise monorepos but requires self-hosted runners for specialized GPU hardware.

Jenkins: Fully open-source and customizable with rich plugin ecosystem; however, requires substantial operational overhead, server provisioning, Groovy scripting, and ongoing security maintenance.

7. Production Deployment Blueprint & Governance

Branch Protection Rules: Enforce branch protection on main requiring all four CI stages to pass and at least one peer approval before merging pull requests.

Zero-Downtime CD Deployments: Successful merges trigger automated Docker image tagging and pushing to GitHub Container Registry (ghcr.io), followed by rolling updates in Kubernetes.

Continuous Training (CT) Triggers: Upstream DVC data changes automatically trigger model retraining workflows, ensuring model freshness without manual developer intervention.

8. Laboratory Review & Viva Voce Reference

Q: Why is CI/CD essential for Machine Learning compared to traditional software?

A: ML systems have three independent vectors of change: code, data, and model parameters. CI/CD automates validation across all three, verifying that code runs, model weights are uncorrupted (DVC), and the containerized API satisfies latency and accuracy SLAs before production deployment.

Q: What is the purpose of the Docker smoke test in Stage 4?

A: Passing local unit tests does not guarantee container runtime success due to OS library discrepancies. The smoke test builds the real Docker image, launches the container, and executes live HTTP requests against /health and /predict to prove runtime viability.

Q: How does DVC integrate with GitHub Actions when large model files are excluded from Git?

A: Small text pointer files (.dvc) containing cryptographic hashes are versioned in Git. During CI execution, the runner uses AWS S3 / GCS credentials stored in GitHub Secrets to execute 'dvc pull', downloading the verified binary model without bloating Git history.

Q: What strategies prevent broken deployments if a model regresses in latency or accuracy?

A: Automated regression gates in Pytest fail the CI pipeline if p95 latency exceeds 150ms or F1 score drops below baseline. In CD, Blue-Green deployments and Canary rollouts ensure traffic is only routed to new containers after automated healthchecks pass.

9. Conclusion

Experiment 7 successfully established an automated, enterprise-grade CI/CD pipeline using GitHub Actions, Pytest, DVC, and Docker. The multi-stage pipeline provides continuous, automated validation—from static linting to live container smoke testing—completing all four stages in 8.31 seconds with a 100% pass rate. This guarantees that only verified, robust, sub-150ms customer support emotion inference microservices are promoted to production.